

Niboon Boonprakob 61340500038

Object Recognition

The objective of this lab is very simple, to recognize objects in images. You will be working with a well-known dataset called CIFAR-10.

You can learn more about this dataset and download it here:

<https://www.cs.toronto.edu/~kriz/cifar.html>

In the webpage above, they also included a few publications based on CIFAR-10 data, which showed some amazing accuracies. The worst network on the page (a shallow convolutional neural network) can classify images with roughly 75% accuracy.

```
In [1]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

1. Write a function to load data

The dataset webpage in the previous section also provide a simple way to load data from your harddrive using pickle. You may use their function for this exercise.

Construct two numpy arrays for train images and train labels from data_batch_1 to data_batch_5. Then, construct two numpy arrays for test images, and test labels from test batch file. The original image size is 32 x 32 x 3. You may flatten the arrays so the final arrays are of size 1 x 3072.

```
In [2]: # Import TensorFlow 2.0
%tensorflow_version 2.x
import tensorflow as tf
import keras
import matplotlib.pyplot as plt
import numpy as np
import random
from tqdm import tqdm

# Check that we are using a GPU, if not switch runtimes
# using Runtime > Change Runtime Type > GPU

assert len(tf.config.list_physical_devices('GPU')) > 0
```

```
In [3]: def unpickle(file):
    import pickle
    with open(file, 'rb') as fo:
        dict = pickle.load(fo, encoding='bytes')
    return dict
```

```
In [4]: pictures=[]
labels = []
for i in range(5):
```

```
file = "/content/drive/MyDrive/Colab Notebooks/Lab 3/dataset/cifar-10-batches-py/d  
# print(file)  
dataset = unpickle(file)  
pictures.append(dataset[b"data"] )  
labels += dataset[b"labels"]  
# print(np.array(pictures).shape)  
pictures = np.concatenate(np.array(pictures) , axis = 0)  
# print(path)
```

In [5]: pictures

```
Out[5]: array([[ 59,  43,  50, ..., 140,  84,  72],  
               [154, 126, 105, ..., 139, 142, 144],  
               [255, 253, 253, ..., 83,  83,  84],  
               ...,  
               [ 35,  40,  42, ...,  77,  66,  50],  
               [189, 186, 185, ..., 169, 171, 171],  
               [229, 236, 234, ..., 173, 162, 161]], dtype=uint8)
```

In [6]: labels

```
Out[6]: [6,  
9,  
9,  
4,  
1,  
1,  
2,  
7,  
8,  
3,  
4,  
7,  
7,  
2,  
9,  
9,  
9,  
3,  
2,  
6,  
4,  
3,  
6,  
6,  
2,  
6,  
3,  
5,  
4,  
0,  
0,  
9,  
1,  
3,  
4,  
0,  
3,  
7,  
3,  
3,  
5,  
2,  
2,  
7,  
1,  
1,  
1,
```

2,
2,
0,
9,
5,
7,
9,
2,
2,
5,
2,
4,
3,
1,
1,
8,
2,
1,
1,
4,
9,
7,
8,
5,
9,
6,
7,
3,
1,
9,
0,
3,
1,
3,
5,
4,
5,
7,
7,
4,
7,
9,
4,
2,
3,
8,
0,
1,
6,
1,
1,
4,
1,
8,
3,
9,
6,
6,
1,
8,
5,
2,
9,
9,
8,
1,
7,
7,
0,

0,
6,
9,
1,
2,
2,
9,
2,
6,
6,
1,
9,
5,
0,
4,
7,
6,
7,
1,
8,
1,
1,
2,
8,
1,
3,
3,
6,
2,
4,
9,
9,
5,
4,
3,
6,
7,
4,
6,
8,
5,
5,
4,
3,
1,
8,
4,
7,
6,
0,
9,
5,
1,
3,
8,
2,
7,
5,
3,
4,
1,
5,
7,
0,
4,
7,
5,
5,
1,

0,
9,
6,
9,
0,
8,
7,
8,
8,
2,
5,
2,
3,
5,
0,
6,
1,
9,
3,
6,
9,
1,
3,
9,
6,
6,
7,
1,
0,
9,
5,
8,
5,
2,
9,
0,
8,
8,
0,
6,
9,
1,
1,
6,
3,
7,
6,
6,
0,
6,
6,
1,
7,
1,
5,
8,
3,
6,
6,
8,
6,
8,
4,
6,
6,
1,
3,
8,
3,

```
4,  
1,  
7,  
1,  
3,  
8,  
5,  
1,  
1,  
4,  
0,  
9,  
3,  
7,  
4,  
9,  
9,  
2,  
4,  
9,  
9,  
1,  
0,  
5,  
9,  
0,  
8,  
2,  
1,  
2,  
0,  
5,  
6,  
3,  
2,  
7,  
8,  
8,  
6,  
0,  
7,  
9,  
4,  
5,  
6,  
4,  
2,  
1,  
1,  
2,  
1,  
5,  
9,  
9,  
0,  
8,  
4,  
1,  
1,  
6,  
3,  
3,  
9,  
0,  
7,  
9,  
7,  
7,  
9,
```

```
1,  
5,  
1,  
6,  
6,  
8,  
7,  
1,  
3,  
0,  
3,  
3,  
2,  
4,  
5,  
7,  
5,  
9,  
0,  
3,  
4,  
0,  
4,  
4,  
6,  
0,  
0,  
6,  
6,  
0,  
8,  
1,  
6,  
2,  
9,  
2,  
5,  
9,  
6,  
7,  
4,  
1,  
8,  
7,  
3,  
6,  
9,  
3,  
0,  
4,  
0,  
5,  
1,  
0,  
3,  
4,  
8,  
5,  
4,  
7,  
2,  
3,  
9,  
7,  
6,  
7,  
1,  
4,  
7,
```

0,
1,
7,
3,
1,
8,
4,
4,
2,
0,
2,
2,
0,
0,
9,
0,
9,
6,
8,
2,
7,
7,
4,
0,
3,
0,
8,
9,
4,
2,
7,
2,
5,
2,
5,
1,
9,
4,
8,
5,
1,
7,
4,
4,
0,
6,
9,
0,
7,
8,
8,
9,
9,
3,
3,
4,
0,
4,
5,
6,
6,
0,
1,
0,
8,
0,
4,
8,
8,

1,
5,
2,
6,
8,
1,
0,
0,
7,
7,
5,
9,
6,
2,
8,
3,
4,
7,
3,
9,
0,
1,
2,
4,
8,
1,
8,
6,
4,
4,
5,
7,
1,
3,
9,
8,
0,
1,
7,
5,
8,
2,
8,
0,
4,
1,
8,
9,
8,
2,
9,
9,
2,
7,
5,
7,
3,
8,
8,
4,
4,
2,
7,
1,
6,
4,
0,
4,
6,

9,
7,
6,
2,
5,
5,
1,
7,
2,
2,
2,
9,
5,
4,
2,
7,
8,
1,
3,
4,
3,
7,
6,
9,
8,
0,
6,
0,
2,
2,
2,
1,
8,
4,
0,
1,
8,
8,
1,
5,
7,
6,
4,
5,
8,
7,
1,
9,
1,
9,
8,
4,
7,
3,
8,
8,
2,
6,
6,
7,
1,
6,
8,
1,
9,
7,
8,
3,
0,

```
1,  
0,  
8,  
8,  
3,  
0,  
0,  
1,  
5,  
0,  
8,  
8,  
7,  
9,  
9,  
0,  
9,  
4,  
1,  
3,  
6,  
6,  
4,  
4,  
7,  
5,  
6,  
0,  
8,  
0,  
3,  
2,  
8,  
4,  
6,  
9,  
9,  
7,  
0,  
3,  
3,  
6,  
7,  
4,  
9,  
1,  
6,  
2,  
7,  
2,  
2,  
0,  
6,  
7,  
5,  
7,  
6,  
8,  
9,  
0,  
9,  
4,  
4,  
7,  
0,  
9,  
4,  
9,  
6,
```

9,
4,
5,
7,
9,
2,
4,
5,
1,
4,
3,
9,
6,
5,
6,
9,
3,
3,
5,
0,
7,
2,
1,
3,
6,
4,
0,
0,
2,
5,
0,
1,
0,
2,
3,
9,
8,
4,
9,
8,
0,
2,
6,
4,
4,
0,
1,
8,
8,
3,
6,
9,
6,
6,
6,
7,
8,
2,
4,
5,
7,
6,
5,
3,
0,
5,
0,
5,
0,
8,

2,
6,
7,
3,
8,
2,
1,
7,
6,
7,
1,
0,
9,
5,
5,
0,
1,
7,
6,
9,
0,
4,
7,
7,
1,
5,
9,
4,
0,
8,
5,
9,
9,
6,
7,
1,
8,
3,
2,
3,
8,
2,
2,
4,
6,
0,
0,
5,
3,
8,
2,
3,
7,
2,
9,
3,
8,
7,
8,
2,
7,
9,
0,
2,
3,
2,
2,
2,
3,

3,
6,
2,
3,
2,
8,
0,
5,
5,
1,
4,
5,
6,
6,
2,
7,
0,
1,
1,
7,
7,
8,
2,
9,
2,
2,
4,
2,
1,
1,
1,
6,
6,
5,
1,
1,
7,
0,
4,
3,
3,
7,
1,
2,
3,
5,
5,
5,
6,
1,
4,
3,
7,
8,
8,
3,
6,
6,
2,
3,
0,
9,
4,
3,
8,
0,
0,
1,
1,

5,
4,
9,
3,
1,
8,
9,
3,
9,
9,
2,
9,
4,
8,
2,
9,
8,
8,
1,
5,
3,
6,
8,
7,
6,
9,
8,
0,
6,
4,
0,
0,
2,
5,
8,
2,
0,
2,
7,
6,
9,
7,
1,
5,
5,
6,
6,
3,
6,
2,
4,
7,
0,
5,
6,
4,
6,
5,
2,
4,
6,
1,
6,
0,
4,
0,
3,
1,
8,

```
5,  
4,  
4,  
1,  
7,  
3,  
9,  
4,  
7,  
9,  
7,  
3,  
7,  
2,  
8,  
4,  
6,  
6,  
1,  
2,  
9,  
0,  
4,  
8,  
7,  
3,  
9,  
8,  
7,  
7,  
0,  
2,  
4,  
1,  
1,  
4,  
1,  
5,  
4,  
0,  
5,  
6,  
2,  
8,  
5,  
0,  
2,  
1,  
3,  
5,  
7,  
3,  
5,  
1,  
3,  
5,  
5,  
...
```

In [7]: `len(pictures) , len(labels)`

Out[7]: (50000, 50000)

In [8]: `pictures.shape`

Out[8]: (50000, 3072)

2. Classify Dogs v.s. Cats

Let's start simple by creating logistic regression model to classify images. We will select only two classes of images for this exercise.

1. From 50,000 train images and 10,000 test images, we want to reduce the data size. Write code to filter only dog images (label = 3) and cat images (label = 5).
2. Create a logistic regression model to classify cats and dogs. Report your accuracy.

```
In [9]: labels.index(3)
```

```
Out[9]: 9
```

```
In [10]: dog_labels = []
dog_img = []
cat_labels = []
cat_img = []

element_dog = 3
element_cat = 5
```

```
for i in range(len(labels)):
    if (labels[i] == element_dog):
        dog_labels.append(labels[i])
        dog_img.append(pictures[i])
    elif (labels[i] == element_cat):
        cat_labels.append(labels[i])
        cat_img.append(pictures[i])
```

```
In [11]: np.array(dog_img).shape , np.array(cat_img).shape
```

```
Out[11]: ((5000, 3072), (5000, 3072))
```

```
In [12]: img_all = [cat_img, dog_img]
# np.array(img_all).shape
d_c_img = np.concatenate(np.array(img_all), axis = 0)
```

```
In [13]: d_c_img.shape
```

```
Out[13]: (10000, 3072)
```

```
In [14]: labels_all = [cat_labels, dog_labels]
d_c_label = np.concatenate(np.array(labels_all), axis=0)
```

```
In [15]: d_c_label.shape
```

```
Out[15]: (10000,)
```

import tensorflow as tf

```
In [16]: test = "/content/drive/MyDrive/Colab Notebooks/Lab 3/dataset/cifar-10-batches-py/test"
y_test = unpickle(test)
y_test_data = y_test[b"data"]
y_test_labels = y_test[b"labels"]
```

```
In [17]: len(y_test_data) , len(y_test_labels)
```

Out[17]: (10000, 10000)

In [18]:

```
# T : images , t:Labels
T_y_test_dog = []
T_y_test_cat = []
t_y_test_dog = []
t_y_test_cat = []
element_dog = 3
element_cat = 5
for i in range(len(y_test_labels)):
    if (y_test_labels[i] == element_cat):
        t_y_test_cat.append(y_test_labels[i])
        T_y_test_cat.append(y_test_data[i])
    elif (y_test_labels[i] == element_dog):
        T_y_test_dog.append(y_test_data[i])
        t_y_test_dog.append(y_test_labels[i])

# y_test dc_img_test : x_test , dc_Labels_test : y_test
img_y_test = [T_y_test_dog , T_y_test_cat]
labels_y_test = [t_y_test_dog , t_y_test_cat]
dc_img_test = np.concatenate(np.array(img_y_test),axis = 0)
dc_labels_test = np.concatenate(np.array(labels_y_test),axis=0)
print(dc_img_test.shape , dc_labels_test.shape)
```

(2000, 3072) (2000,)

In [18]:

In [19]:

```
from sklearn.linear_model import LogisticRegression
clf = LogisticRegression(random_state=42,max_iter=1000).fit(dc_img, dc_label)
y_pred = clf.predict(dc_img_test)
```

/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py:940: ConvergenceWarning: lbfsgs failed to converge (status=1):
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)

In [20]:

```
from sklearn.metrics import accuracy_score
print(f"Accuracy: {accuracy_score(dc_labels_test , y_pred)}")
```

Accuracy: 0.561

3. The Real Challenge

The majority of your score for this lab will come from this real challenge. You are going to construct a neural network model to classify 10 classes of images from CIFAR-10 dataset. You will get half the credits for this one if you complete the assignment, and will get another half if you can exceed the target accuracy of 75%. (You may use any combination of sklearn, opencv, or tensorflow to do this exercise).

Design at least 3 variants of neural network models. Each model should have different architectures. (Do not vary just a few parameters, the architecture of the network must change in each model). In your notebook, explain your experiments in details and display the accuracy score for each experiment.

```
In [21]: ### Load dataset ###  
  
# X_train : pictures  
# y_train : Labels  
# X_test : y_test_data  
# y_test : y_test_labels
```

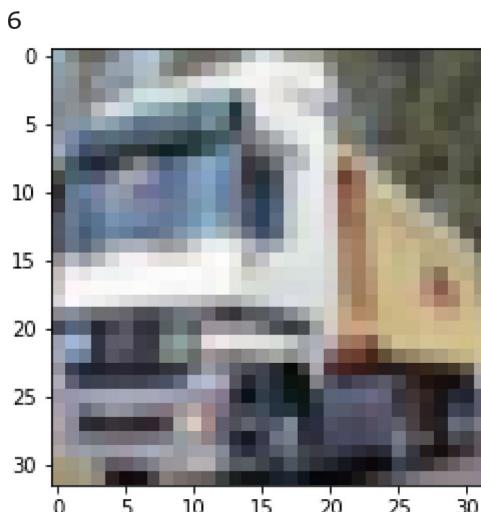
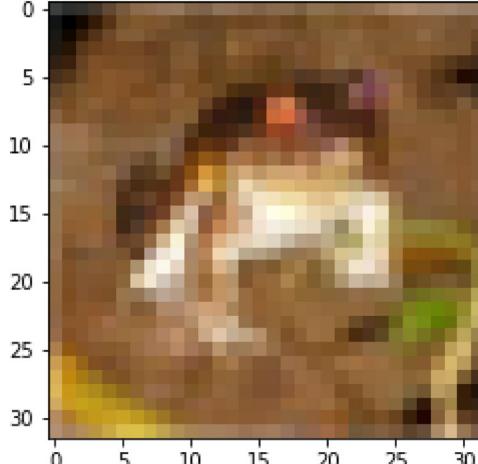
```
In [22]: pictures.shape
```

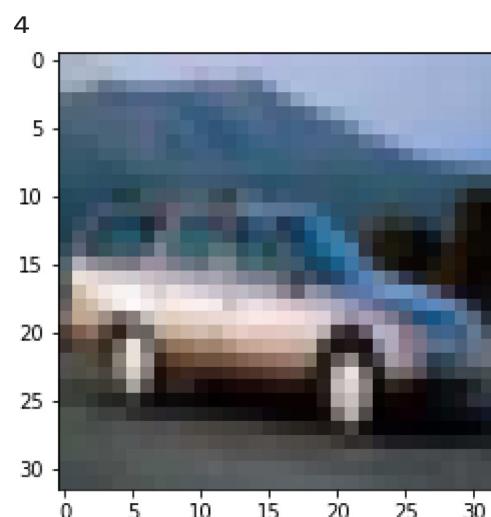
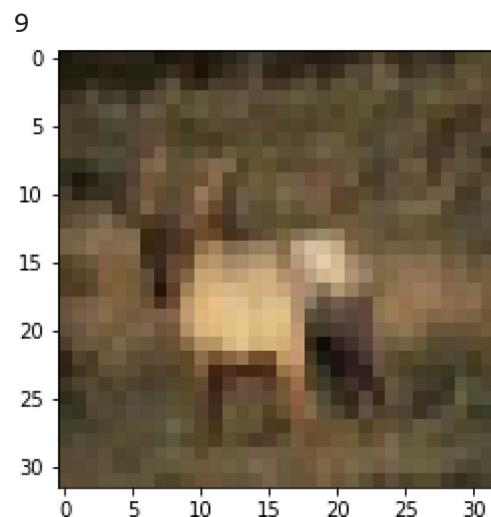
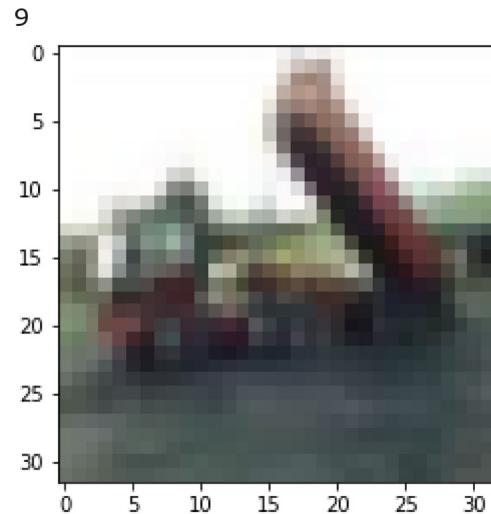
```
Out[22]: (50000, 3072)
```

```
In [23]: X_train = []  
X_test = []  
for i in range(len(pictures)):  
    img = np.reshape(pictures[i], (3,32,32))  
    X_train.append(np.transpose(img, (1, 2, 0)))  
    print(X_train[0].shape)  
  
for i in range(len(y_test_data)):  
    img = np.reshape(y_test_data[i], (3,32,32))  
    X_test.append(np.transpose(img, (1, 2, 0)))  
y_train = labels  
y_test = y_test_labels
```

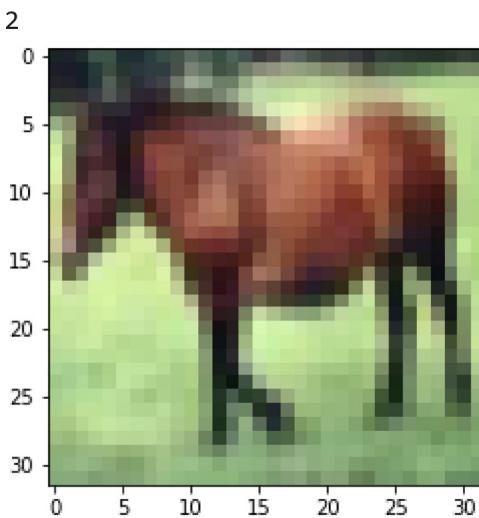
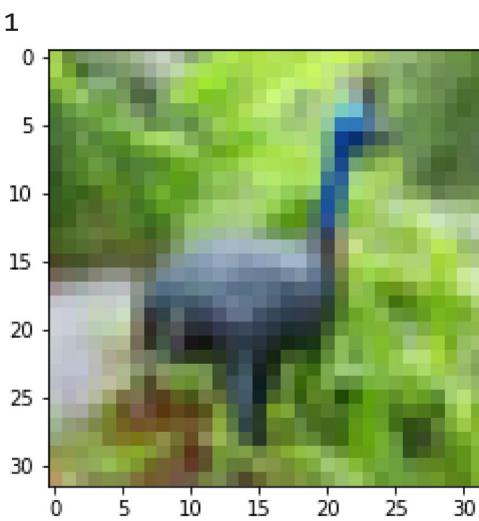
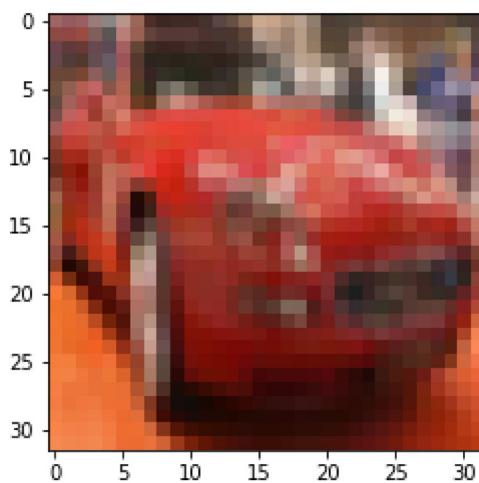
```
(32, 32, 3)
```

```
In [24]: import matplotlib.pyplot as plt  
for i in range(10):  
    plt.imshow(X_train[i])  
    plt.show()  
    print(y_train[i])
```

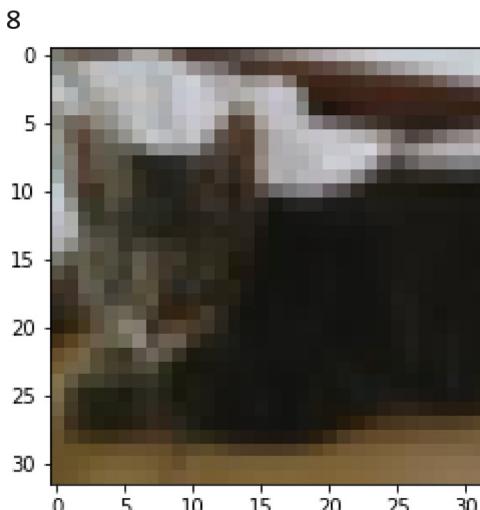
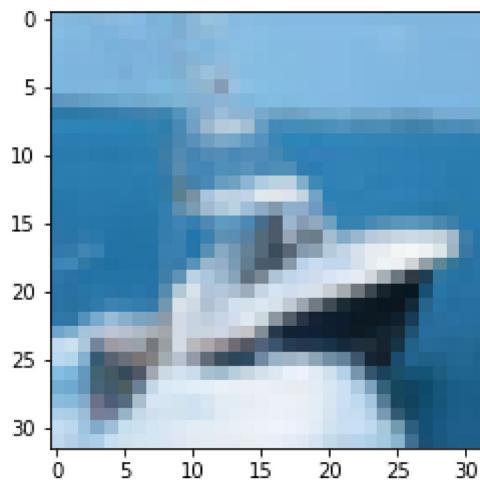




1

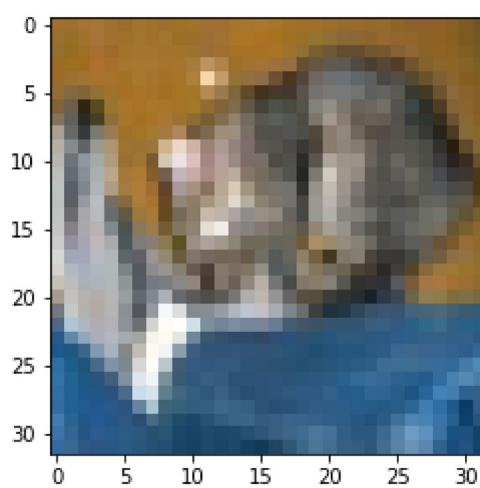


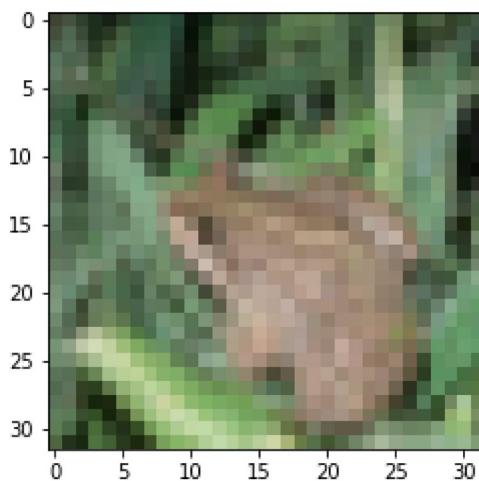
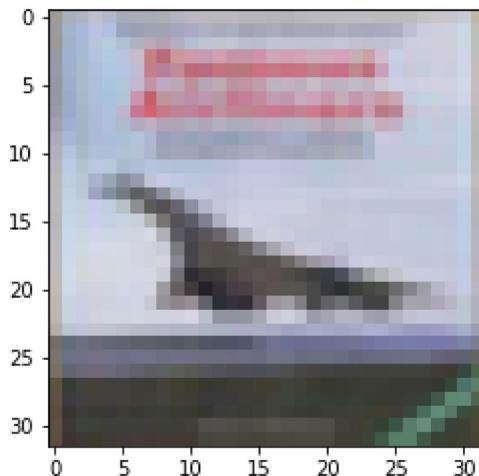
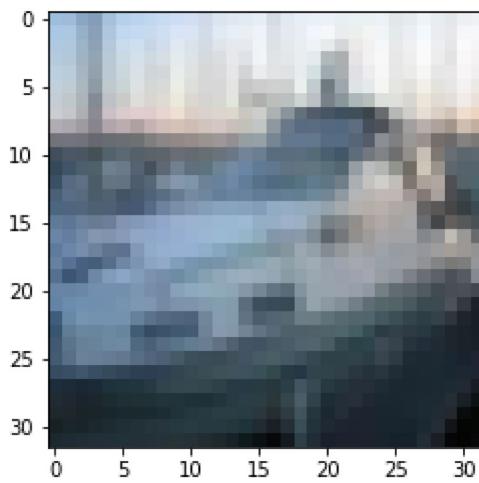
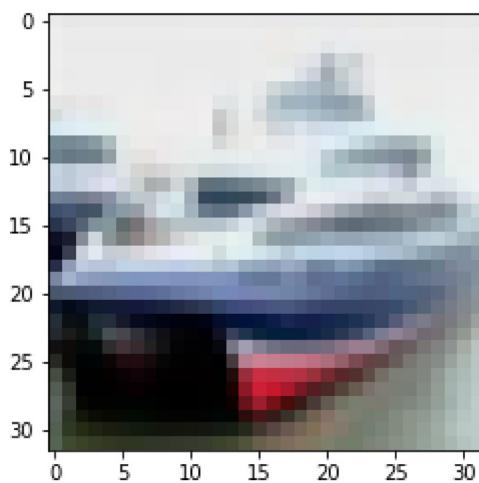
7

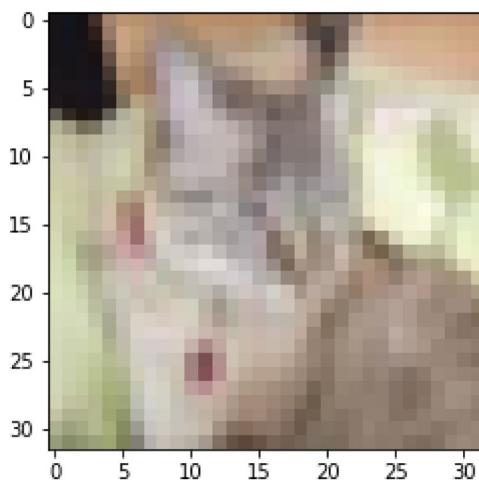
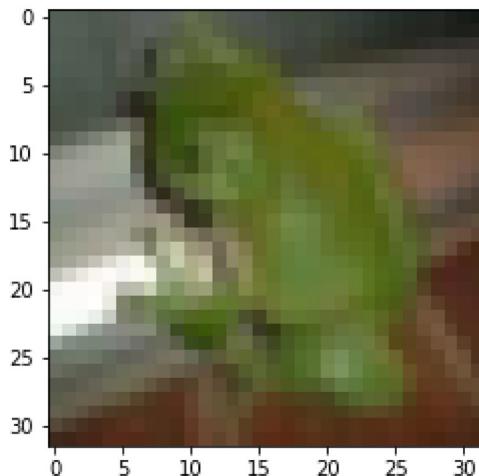
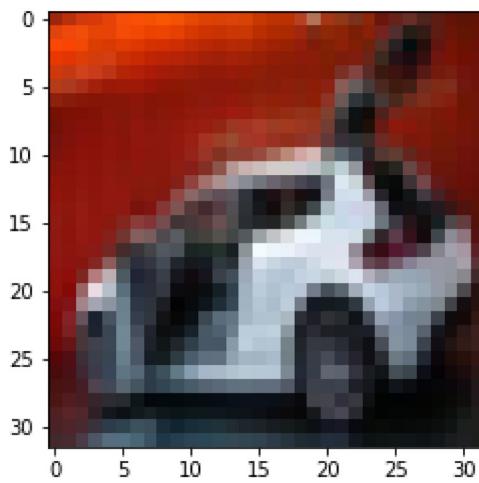
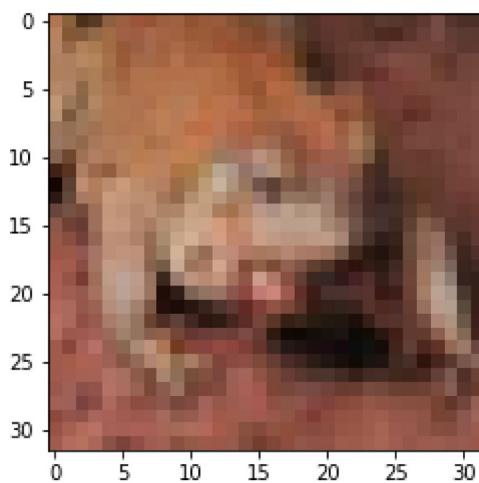


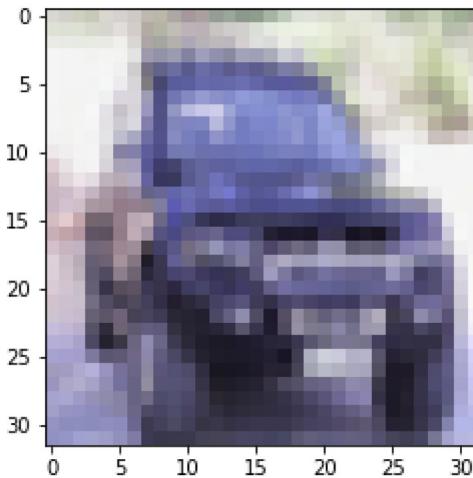
3

```
In [25]: for i in range(10):
    plt.imshow(X_test[i])
    plt.show()
```









In [26]: # 0->1

```
train_images = (np.expand_dims(X_train, axis=-1)/255.).astype(np.float32)
# train_labels = (y_train).astype(np.int64)
test_images = (np.expand_dims(X_test, axis=-1)/255.).astype(np.float32)
# test_labels = (y_test ).astype(np.int64)
```

In [27]: train_images = train_images.reshape(train_images.shape[:-1])
print(train_images.shape)

(50000, 32, 32, 3)

In [28]: test_images= test_images.reshape(test_images.shape[:-1])
print(test_images.shape)

(10000, 32, 32, 3)

In [29]: `def define_model():
 model = tf.keras.Sequential()
 model.add(tf.keras.layers.Convolution2D(32, (3, 3), activation='relu', kernel_initializer='he_normal'))
 model.add(tf.keras.layers.Convolution2D(32, (3, 3), activation='relu', kernel_initializer='he_normal'))
 model.add(tf.keras.layers.MaxPooling2D((2, 2)))
 model.add(tf.keras.layers.Flatten())
 model.add(tf.keras.layers.Dense(128, activation='relu', kernel_initializer='he_normal'))
 model.add(tf.keras.layers.Dense(10, activation='softmax'))
 # compile model
 opt = tf.keras.optimizers.SGD(lr=0.001, momentum=0.9)
 model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])
 return model`

```
/usr/local/lib/python3.7/dist-packages/tensorflow/python/keras/optimizer_v2/optimizer_v2.py:375: UserWarning: The `lr` argument is deprecated, use `learning_rate` instead.
    "The `lr` argument is deprecated, use `learning_rate` instead.")
```

In [30]: `from sklearn import preprocessing
lb = preprocessing.LabelBinarizer()
lb.fit(labels)
y_train = lb.transform(labels)
y_test = lb.transform(y_test_labels)`

In [31]: print(y_train)

```
[[0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 1]
 [0 0 0 ... 0 0 1]
 ...
 [0 0 0 ... 0 0 1]]
```

```
[0 1 0 ... 0 0 0]
[0 1 0 ... 0 0 0]]
```

In [32]: `print(y_test)`

```
[[0 0 0 ... 0 0 0]
[0 0 0 ... 0 1 0]
[0 0 0 ... 0 1 0]
...
[0 0 0 ... 0 0 0]
[0 1 0 ... 0 0 0]
[0 0 0 ... 1 0 0]]
```

In [33]: `print(train_images)`

```
[[[[0.23137255 0.24313726 0.24705882]
[0.16862746 0.18039216 0.1764706 ]
[0.19607843 0.1882353 0.16862746]
...
[0.61960787 0.5176471 0.42352942]
[0.59607846 0.49019608 0.4
[0.5803922 0.4862745 0.40392157]]
[[0.0627451 0.07843138 0.07843138]
[0. 0. 0. ]
[0.07058824 0.03137255 0.
...
[0.48235294 0.34509805 0.21568628]
[0.46666667 0.3254902 0.19607843]
[0.47843137 0.34117648 0.22352941]]
[[0.09803922 0.09411765 0.08235294]
[0.0627451 0.02745098 0.
[0.19215687 0.10588235 0.03137255]
...
[0.4627451 0.32941177 0.19607843]
[0.47058824 0.32941177 0.19607843]
[0.42745098 0.28627452 0.16470589]]
...
[[0.8156863 0.6666667 0.3764706 ]
[0.7882353 0.6 0.13333334]
[0.7764706 0.6313726 0.10196079]
...
[0.627451 0.52156866 0.27450982]
[0.21960784 0.12156863 0.02745098]
[0.20784314 0.13333334 0.07843138]]
[[0.7058824 0.54509807 0.3764706 ]
[0.6784314 0.48235294 0.16470589]
[0.7294118 0.5647059 0.11764706]
...
[0.72156864 0.5803922 0.36862746]
[0.38039216 0.24313726 0.13333334]
[0.3254902 0.20784314 0.13333334]]
[[0.69411767 0.5647059 0.45490196]
[0.65882355 0.5058824 0.36862746]
[0.7019608 0.5568628 0.34117648]
...
[0.84705883 0.72156864 0.54901963]
[0.5921569 0.4627451 0.32941177]
[0.48235294 0.36078432 0.28235295]]]
[[[0.6039216 0.69411767 0.73333335]
[0.49411765 0.5372549 0.53333336]
[0.4117647 0.40784314 0.37254903]]]
```

```

...
[[0.35686275 0.37254903 0.2784314 ]
[0.34117648 0.3529412 0.2784314 ]
[0.30980393 0.31764707 0.27450982]]]

[[[0.54901963 0.627451 0.6627451 ]
[0.5686275 0.6 0.6039216 ]
[0.49019608 0.49019608 0.4627451 ]
...
[0.3764706 0.3882353 0.30588236]
[0.3019608 0.3137255 0.24313726]
[0.2784314 0.28627452 0.23921569]]]

[[[0.54901963 0.60784316 0.6431373 ]
[0.54509807 0.57254905 0.58431375]
[0.4509804 0.4509804 0.4392157 ]
...
[0.30980393 0.32156864 0.2509804 ]
[0.26666668 0.27450982 0.21568628]
[0.2627451 0.27058825 0.21568628]]]

...
[[[0.6862745 0.654902 0.6509804 ]
[0.6117647 0.6039216 0.627451 ]
[0.6039216 0.627451 0.6666667 ]
...
[0.16470589 0.13333334 0.14117648]
[0.23921569 0.20784314 0.22352941]
[0.3647059 0.3254902 0.35686275]]]

[[[0.64705884 0.6039216 0.5019608 ]
[0.6117647 0.59607846 0.50980395]
[0.62352943 0.6313726 0.5568628 ]
...
[0.40392157 0.3647059 0.3764706 ]
[0.48235294 0.44705883 0.47058824]
[0.5137255 0.4745098 0.5137255 ]]

[[[0.6392157 0.5803922 0.47058824]
[0.61960787 0.5803922 0.47843137]
[0.6392157 0.6117647 0.52156866]
...
[0.56078434 0.52156866 0.54509807]
[0.56078434 0.5254902 0.5568628 ]
[0.56078434 0.52156866 0.5647059 ]]]]

[[[[1. 1. 1.
[0.99215686 0.99215686 0.99215686]
[0.99215686 0.99215686 0.99215686]
...
[0.99215686 0.99215686 0.99215686]
[0.99215686 0.99215686 0.99215686]
[0.99215686 0.99215686 0.99215686]]]

[[[1. 1. 1.
[1. 1. 1.
[1. 1. 1.
...
[1. 1. 1.
[1. 1. 1.
[1. 1. 1.]]

[[[1. 1. 1.
[0.99607843 0.99607843 0.99607843]
[0.99607843 0.99607843 0.99607843]
...
[0.99607843 0.99607843 0.99607843]
```

```

[[0.99607843 0.99607843 0.99607843]
 [0.99607843 0.99607843 0.99607843]]

...
[[0.44313726 0.47058824 0.4392157 ]
 [0.43529412 0.4627451 0.43529412]
 [0.4117647 0.4392157 0.41568628]
 ...
[[0.28235295 0.31764707 0.3137255 ]
 [0.28235295 0.3137255 0.30980393]
 [0.28235295 0.3137255 0.30980393]]

[[0.43529412 0.4627451 0.43137255]
 [0.40784314 0.43529412 0.40784314]
 [0.3882353 0.41568628 0.38431373]
 ...
[[0.26666668 0.29411766 0.28627452]
 [0.27450982 0.29803923 0.29411766]
 [0.30588236 0.32941177 0.32156864]]

[[0.41568628 0.44313726 0.4117647 ]
 [0.3882353 0.41568628 0.38431373]
 [0.37254903 0.4 0.36862746]
 ...
[[0.30588236 0.33333334 0.3254902 ]
 [0.30980393 0.33333334 0.3254902 ]
 [0.3137255 0.3372549 0.32941177]]]

...
[[[0.13725491 0.69803923 0.92156863]
 [0.15686275 0.6901961 0.9372549]
 [0.16470589 0.6901961 0.94509804]
 ...
[[0.3882353 0.69411767 0.85882354]
 [0.30980393 0.5764706 0.77254903]
 [0.34901962 0.5803922 0.7411765 ]]

[[0.22352941 0.7137255 0.91764706]
 [0.17254902 0.72156864 0.98039216]
 [0.19607843 0.7176471 0.9411765 ]
 ...
[[0.6117647 0.7137255 0.78431374]
 [0.5529412 0.69411767 0.80784315]
 [0.45490196 0.58431375 0.6862745 ]]

[[0.38431373 0.77254903 0.92941177]
 [0.2509804 0.7411765 0.9882353]
 [0.27058825 0.7529412 0.9607843 ]
 ...
[[0.7372549 0.7647059 0.80784315]
 [0.46666667 0.5294118 0.5764706 ]
 [0.23921569 0.30980393 0.3529412 ]]

...
[[0.28627452 0.30980393 0.3019608 ]
 [0.20784314 0.24705882 0.26666668]
 [0.21176471 0.26666668 0.3137255 ]
 ...
[[0.06666667 0.15686275 0.2509804 ]
 [0.08235294 0.14117648 0.2
 [0.12941177 0.1882353 0.19215687]]]

[[0.23921569 0.26666668 0.29411766]
 [0.21568628 0.27450982 0.3372549 ]]
```

```
[0.22352941 0.30980393 0.40392157]
...
[0.09411765 0.1882353 0.28235295]
[0.06666667 0.13725491 0.20784314]
[0.02745098 0.09019608 0.1254902 ]]

[[[0.17254902 0.21960784 0.28627452]
[0.18039216 0.25882354 0.34509805]
[0.19215687 0.3019608 0.4117647 ]
...
[0.10588235 0.20392157 0.3019608 ]
[0.08235294 0.16862746 0.25882354]
[0.04705882 0.12156863 0.19607843]]]

[[[0.7411765 0.827451 0.9411765 ]
[0.7294118 0.8156863 0.9254902 ]
[0.7254902 0.8117647 0.92156863]
...
[0.6862745 0.7647059 0.8784314 ]
[0.6745098 0.7607843 0.87058824]
[0.6627451 0.7607843 0.8627451 ]]

[[0.7607843 0.8235294 0.9372549 ]
[0.7490196 0.8117647 0.9254902 ]
[0.74509805 0.80784315 0.92156863]
...
[0.6784314 0.7529412 0.8627451 ]
[0.67058825 0.7490196 0.85490197]
[0.654902 0.74509805 0.84705883]]]

[[0.8156863 0.85882354 0.95686275]
[0.8039216 0.84705883 0.9411765 ]
[0.8 0.84313726 0.9372549 ]
...
[0.6862745 0.7490196 0.8509804 ]
[0.6745098 0.74509805 0.84705883]
[0.6627451 0.7490196 0.84313726]]]

...
[[0.8117647 0.78039217 0.70980394]
[0.79607844 0.7647059 0.6862745 ]
[0.79607844 0.76862746 0.6784314 ]
...
[0.5294118 0.5176471 0.49803922]
[0.63529414 0.61960787 0.5882353 ]
[0.65882355 0.6392157 0.5921569 ]]

[[0.7764706 0.74509805 0.6666667 ]
[0.7411765 0.70980394 0.62352943]
[0.7058824 0.6745098 0.5764706 ]
...
[0.69803923 0.67058825 0.627451 ]
[0.6862745 0.6627451 0.6117647 ]
[0.6862745 0.6627451 0.6039216 ]]

[[0.7764706 0.7411765 0.6784314 ]
[0.7411765 0.70980394 0.63529414]
[0.69803923 0.6666667 0.58431375]
...
[0.7647059 0.72156864 0.6627451 ]
[0.76862746 0.7411765 0.67058825]
[0.7647059 0.74509805 0.67058825]]]

[[[0.8980392 0.8980392 0.9372549 ]
[0.9254902 0.92941177 0.96862745]
[0.91764706 0.9254902 0.96862745]
```

```

...
[[0.8509804  0.85882354 0.9137255 ]
 [0.8666667  0.8745098  0.91764706]
 [0.87058824 0.8745098  0.9137255 ]]

[[[0.87058824 0.8666667 0.8980392 ]
 [0.9372549  0.9372549  0.9764706]
 [0.9137255  0.91764706 0.9647059]
 ...
 [[0.8745098  0.8745098  0.9254902 ]
 [0.8901961  0.89411765 0.93333334]
 [0.8235294  0.827451   0.8627451 ]]

[[[0.8352941  0.80784315 0.827451 ]
 [0.91764706 0.9098039  0.9372549]
 [0.90588236 0.9137255  0.95686275]
 ...
 [[0.8627451  0.8627451  0.9098039 ]
 [0.8627451  0.85882354 0.9098039 ]
 [0.7921569  0.79607844 0.84313726]]

...
[[[0.5882353  0.56078434 0.5294118 ]
 [0.54901963 0.5294118  0.49803922]
 [0.5176471  0.49803922 0.47058824]
 ...
 [[0.8784314  0.87058824 0.85490197]
 [0.9019608  0.89411765 0.88235295]
 [0.94509804 0.94509804 0.93333334]]]

[[[0.5372549  0.5176471  0.49411765]
 [0.50980395 0.49803922 0.47058824]
 [0.49019608 0.4745098  0.4509804 ]
 ...
 [[0.70980394 0.7058824  0.69803923]
 [0.7921569  0.7882353  0.7764706 ]
 [0.83137256 0.827451   0.8117647 ]]

[[[0.47843137 0.46666667 0.44705883]
 [0.4627451  0.45490196 0.43137255]
 [0.47058824 0.45490196 0.43529412]
 ...
 [[0.7019608  0.69411767 0.6784314 ]
 [0.6431373  0.6431373  0.63529414]
 [0.6392157  0.6392157  0.6313726 ]]]]

```

In [34]: `test_images.shape`

Out[34]: `(10000, 32, 32, 3)`

In [35]: `y_test.shape`

Out[35]: `(10000, 10)`

Variable

- `X_train` : `train_images`
- `y_train` : `np.array(y_train)`
- `X_test` : `test_images`
- `y_test` : `y_test`

In [36]: `# Define the batch size and the number of epochs to use during training
BATCH_SIZE = 100 #`

```
EPOCHS = 40

'''TODO: Use model.fit to train the CNN model, with the same batch_size and number o
H = model.fit(train_images,np.array(y_train), validation_data=(test_images , y_test)

Epoch 1/40
500/500 [=====] - 3s 5ms/step - loss: 1.8241 - accuracy: 0.
3541 - val_loss: 1.6231 - val_accuracy: 0.4260
Epoch 2/40
500/500 [=====] - 2s 4ms/step - loss: 1.4876 - accuracy: 0.
4704 - val_loss: 1.4034 - val_accuracy: 0.4980
Epoch 3/40
500/500 [=====] - 2s 4ms/step - loss: 1.3434 - accuracy: 0.
5236 - val_loss: 1.3056 - val_accuracy: 0.5352
Epoch 4/40
500/500 [=====] - 2s 4ms/step - loss: 1.2596 - accuracy: 0.
5570 - val_loss: 1.2358 - val_accuracy: 0.5609
Epoch 5/40
500/500 [=====] - 2s 4ms/step - loss: 1.1949 - accuracy: 0.
5819 - val_loss: 1.1951 - val_accuracy: 0.5792
Epoch 6/40
500/500 [=====] - 2s 4ms/step - loss: 1.1379 - accuracy: 0.
6023 - val_loss: 1.1488 - val_accuracy: 0.5934
Epoch 7/40
500/500 [=====] - 2s 4ms/step - loss: 1.0909 - accuracy: 0.
6173 - val_loss: 1.1368 - val_accuracy: 0.5979
Epoch 8/40
500/500 [=====] - 2s 4ms/step - loss: 1.0409 - accuracy: 0.
6372 - val_loss: 1.1131 - val_accuracy: 0.6038
Epoch 9/40
500/500 [=====] - 2s 4ms/step - loss: 1.0010 - accuracy: 0.
6509 - val_loss: 1.0758 - val_accuracy: 0.6210
Epoch 10/40
500/500 [=====] - 2s 4ms/step - loss: 0.9599 - accuracy: 0.
6656 - val_loss: 1.0502 - val_accuracy: 0.6350
Epoch 11/40
500/500 [=====] - 2s 4ms/step - loss: 0.9253 - accuracy: 0.
6779 - val_loss: 1.0476 - val_accuracy: 0.6320
Epoch 12/40
500/500 [=====] - 2s 4ms/step - loss: 0.8821 - accuracy: 0.
6935 - val_loss: 1.0280 - val_accuracy: 0.6428
Epoch 13/40
500/500 [=====] - 2s 4ms/step - loss: 0.8483 - accuracy: 0.
7064 - val_loss: 0.9994 - val_accuracy: 0.6517
Epoch 14/40
500/500 [=====] - 2s 4ms/step - loss: 0.8103 - accuracy: 0.
7199 - val_loss: 1.0018 - val_accuracy: 0.6518
Epoch 15/40
500/500 [=====] - 2s 4ms/step - loss: 0.7788 - accuracy: 0.
7321 - val_loss: 1.0101 - val_accuracy: 0.6517
Epoch 16/40
500/500 [=====] - 2s 4ms/step - loss: 0.7508 - accuracy: 0.
7397 - val_loss: 0.9822 - val_accuracy: 0.6624
Epoch 17/40
500/500 [=====] - 2s 4ms/step - loss: 0.7111 - accuracy: 0.
7553 - val_loss: 0.9730 - val_accuracy: 0.6678
Epoch 18/40
500/500 [=====] - 2s 4ms/step - loss: 0.6863 - accuracy: 0.
7648 - val_loss: 0.9715 - val_accuracy: 0.6707
Epoch 19/40
500/500 [=====] - 2s 4ms/step - loss: 0.6481 - accuracy: 0.
7772 - val_loss: 0.9816 - val_accuracy: 0.6681
Epoch 20/40
500/500 [=====] - 2s 4ms/step - loss: 0.6203 - accuracy: 0.
7881 - val_loss: 1.0344 - val_accuracy: 0.6583
Epoch 21/40
500/500 [=====] - 2s 4ms/step - loss: 0.5873 - accuracy: 0.
7996 - val_loss: 0.9935 - val_accuracy: 0.6690
Epoch 22/40
```

```

500/500 [=====] - 2s 4ms/step - loss: 0.5597 - accuracy: 0.
8107 - val_loss: 0.9906 - val_accuracy: 0.6756
Epoch 23/40
500/500 [=====] - 2s 4ms/step - loss: 0.5290 - accuracy: 0.
8211 - val_loss: 1.0260 - val_accuracy: 0.6678
Epoch 24/40
500/500 [=====] - 2s 4ms/step - loss: 0.4953 - accuracy: 0.
8353 - val_loss: 1.0491 - val_accuracy: 0.6690
Epoch 25/40
500/500 [=====] - 2s 4ms/step - loss: 0.4727 - accuracy: 0.
8422 - val_loss: 1.0432 - val_accuracy: 0.6710
Epoch 26/40
500/500 [=====] - 2s 4ms/step - loss: 0.4403 - accuracy: 0.
8551 - val_loss: 1.0740 - val_accuracy: 0.6645
Epoch 27/40
500/500 [=====] - 2s 4ms/step - loss: 0.4133 - accuracy: 0.
8657 - val_loss: 1.0871 - val_accuracy: 0.6633
Epoch 28/40
500/500 [=====] - 2s 4ms/step - loss: 0.3806 - accuracy: 0.
8772 - val_loss: 1.1226 - val_accuracy: 0.6617
Epoch 29/40
500/500 [=====] - 2s 4ms/step - loss: 0.3525 - accuracy: 0.
8861 - val_loss: 1.1696 - val_accuracy: 0.6600
Epoch 30/40
500/500 [=====] - 2s 4ms/step - loss: 0.3272 - accuracy: 0.
8974 - val_loss: 1.1639 - val_accuracy: 0.6665
Epoch 31/40
500/500 [=====] - 2s 5ms/step - loss: 0.2977 - accuracy: 0.
9078 - val_loss: 1.2233 - val_accuracy: 0.6606
Epoch 32/40
500/500 [=====] - 2s 4ms/step - loss: 0.2705 - accuracy: 0.
9176 - val_loss: 1.2534 - val_accuracy: 0.6573
Epoch 33/40
500/500 [=====] - 2s 4ms/step - loss: 0.2477 - accuracy: 0.
9260 - val_loss: 1.2951 - val_accuracy: 0.6588
Epoch 34/40
500/500 [=====] - 2s 4ms/step - loss: 0.2276 - accuracy: 0.
9319 - val_loss: 1.3184 - val_accuracy: 0.6623
Epoch 35/40
500/500 [=====] - 2s 4ms/step - loss: 0.1988 - accuracy: 0.
9438 - val_loss: 1.3586 - val_accuracy: 0.6609
Epoch 36/40
500/500 [=====] - 2s 4ms/step - loss: 0.1770 - accuracy: 0.
9516 - val_loss: 1.3908 - val_accuracy: 0.6621
Epoch 37/40
500/500 [=====] - 2s 4ms/step - loss: 0.1581 - accuracy: 0.
9578 - val_loss: 1.4702 - val_accuracy: 0.6539
Epoch 38/40
500/500 [=====] - 2s 4ms/step - loss: 0.1425 - accuracy: 0.
9630 - val_loss: 1.4726 - val_accuracy: 0.6580
Epoch 39/40
500/500 [=====] - 2s 4ms/step - loss: 0.1224 - accuracy: 0.
9711 - val_loss: 1.5212 - val_accuracy: 0.6575
Epoch 40/40
500/500 [=====] - 2s 4ms/step - loss: 0.1077 - accuracy: 0.
9761 - val_loss: 1.5499 - val_accuracy: 0.6596

```

Evaluate accuracy on the test dataset

```
In [37]: test_loss, test_acc = model.evaluate(test_images , y_test)# TODO
print('Test accuracy:', test_acc)
```

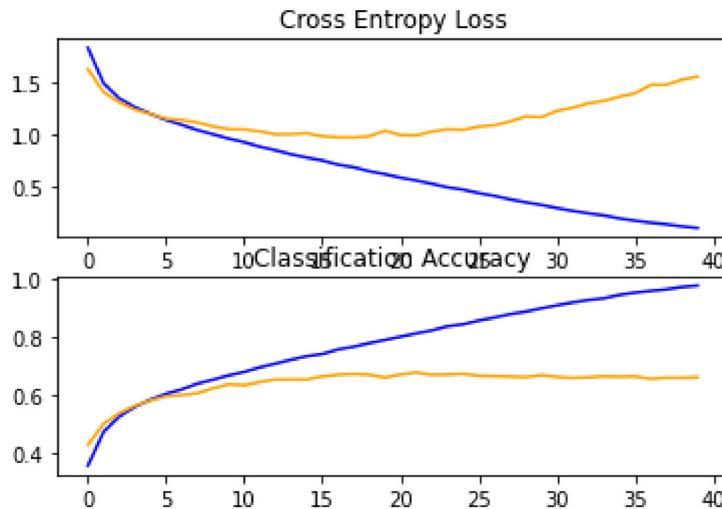
```
313/313 [=====] - 1s 3ms/step - loss: 1.5499 - accuracy: 0.
6596
Test accuracy: 0.659600019454956
```

```
In [38]: import sys
def summarize_diagnostics(history):
    # plot loss
    plt.subplot(211)
    plt.title('Cross Entropy Loss')
    plt.plot(history.history['loss'], color='blue', label='train')
    plt.plot(history.history['val_loss'], color='orange', label='test')
    # plot accuracy
    plt.subplot(212)
    plt.title('Classification Accuracy')
    plt.plot(history.history['accuracy'], color='blue', label='train')
    plt.plot(history.history['val_accuracy'], color='orange', label='test')
    plt.show()
```

```
In [39]: print(H.history)

{'loss': [1.8241266012191772, 1.4876279830932617, 1.343402624130249, 1.2595564126968384, 1.1949219703674316, 1.1378787755966187, 1.0909010171890259, 1.0409411191940308, 1.001021146774292, 0.9599418640136719, 0.925335705280304, 0.8820815086364746, 0.8483354449272156, 0.810336709022522, 0.7787525653839111, 0.7507606148719788, 0.7111119031906128, 0.6862810254096985, 0.6480662822723389, 0.6203339099884033, 0.5872994661331177, 0.559696614742279, 0.5289755463600159, 0.4953438937664032, 0.4726831614971161, 0.44025900959968567, 0.41333696246147156, 0.38063371181488037, 0.35253390669822693, 0.32717305421829224, 0.2977215647697449, 0.270526260137558, 0.2476857602596283, 0.227588951587677, 0.19875311851501465, 0.1770431250333786, 0.1581188440322876, 0.14254115521907806, 0.12238401174545288, 0.10766598582267761], 'accuracy': [0.35405999422073364, 0.4703800082206726, 0.523639976978302, 0.5569999814033508, 0.5819000005722046, 0.6022599935531616, 0.6173200011253357, 0.6372399926185608, 0.6509400010108948, 0.6656200289726257, 0.6778600215911865, 0.6935399770736694, 0.7063800096511841, 0.7198600172996521, 0.7320799827575684, 0.7397400140762329, 0.7553399801254272, 0.7648199796676636, 0.7771999835968018, 0.7880799770355225, 0.7996399998664856, 0.8107200264930725, 0.8210999965667725, 0.835319995880127, 0.8422200083732605, 0.855139970779419, 0.8657000064849854, 0.8771600127220154, 0.8860999941825867, 0.897379994392395, 0.9078400135040283, 0.917639970779419, 0.9259799718856812, 0.9319000244140625, 0.9437999725341797, 0.9515799880027771, 0.9577599763870239, 0.962999995231628, 0.9711199998855591, 0.9760599732398987], 'val_loss': [1.623051643371582, 1.4033665657043457, 1.3055610656738281, 1.2357871532440186, 1.1950500011444092, 1.1488419771194458, 1.1368329524993896, 1.1130682229995728, 1.0757778882980347, 1.0502017736434937, 1.0475515127182007, 1.027968406677246, 0.9993643760681152, 1.0018354654312134, 1.0100650787353516, 0.9822205901145935, 0.9730322957038879, 0.9715232253074646, 0.9816243052482605, 1.0343992710113525, 0.9935475587844849, 0.9906429648399353, 1.0260231494903564, 1.0490942001342773, 1.0431628227233887, 1.074043869972229, 1.0871212482452393, 1.1226484775543213, 1.1696045398712158, 1.1638693809509277, 1.223250150680542, 1.2534306049346924, 1.2951337099075317, 1.3184103965759277, 1.3586204051971436, 1.390756368637085, 1.4702092409133911, 1.4726176261901855, 1.5211753845214844, 1.549892783164978], 'val_accuracy': [0.4259999990463257, 0.49799999594688416, 0.5351999998092651, 0.5608999729156494, 0.579200029373169, 0.5934000015258789, 0.5978999733924866, 0.6037999987602234, 0.6209999918937683, 0.6349999904632568, 0.6320000290870667, 0.642799973487854, 0.6517000198364258, 0.6517999768257141, 0.6517000198364258, 0.6624000072479248, 0.667800092506409, 0.6707000136375427, 0.668099994277954, 0.65829998254776, 0.6690000295639038, 0.6710000038146973, 0.664499980926514, 0.6632999777793884, 0.6617000102996826, 0.6600000262260437, 0.6664999723434448, 0.6606000065803528, 0.6572999954223633, 0.6588000059127808, 0.6622999906539917, 0.6608999967575073, 0.6621000170707703, 0.6539000272750854, 0.657999923706055, 0.6575000286102295, 0.659600019454956]}
```

```
In [40]: summarize_diagnostics(H)
```



reference

<https://machinelearningmastery.com/how-to-develop-a-cnn-from-scratch-for-cifar-10-photo-classification/>

```
In [41]: # define cnn model
def define_model2():
    model = tf.keras.Sequential()
    model.add(tf.keras.layers.Convolution2D(32, (3, 3), activation='relu', kernel_initializer='he_normal'))
    model.add(tf.keras.layers.Convolution2D(32, (3, 3), activation='relu', kernel_initializer='he_normal'))
    model.add(tf.keras.layers.MaxPooling2D((2, 2)))
    model.add(tf.keras.layers.Dropout(0.2))
    model.add(tf.keras.layers.Convolution2D(64, (3, 3), activation='relu', kernel_initializer='he_normal'))
    model.add(tf.keras.layers.Convolution2D(64, (3, 3), activation='relu', kernel_initializer='he_normal'))
    model.add(tf.keras.layers.MaxPooling2D((2, 2)))
    model.add(tf.keras.layers.Dropout(0.2))
    model.add(tf.keras.layers.Convolution2D(128, (3, 3), activation='relu', kernel_initializer='he_normal'))
    model.add(tf.keras.layers.Convolution2D(128, (3, 3), activation='relu', kernel_initializer='he_normal'))
    model.add(tf.keras.layers.MaxPooling2D((2, 2)))
    model.add(tf.keras.layers.Dropout(0.2))
    model.add(tf.keras.layers.Flatten())
    model.add(tf.keras.layers.Dense(128, activation='relu', kernel_initializer='he_normal'))
    model.add(tf.keras.layers.Dropout(0.2))
    model.add(tf.keras.layers.Dense(10, activation='softmax'))
    # compile model
    opt = tf.keras.optimizers.SGD(lr=0.001, momentum=0.9)
    model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])
    return model
```

```
In [42]: model2 = define_model2()
```

```
/usr/local/lib/python3.7/dist-packages/tensorflow/python/keras/optimizer_v2/optimizer_v2.py:375: UserWarning: The `lr` argument is deprecated, use `learning_rate` instead.
  "The `lr` argument is deprecated, use `learning_rate` instead.")
```

```
In [43]: BATCH_SIZE = 100 #
EPOCHS = 40

'''TODO: Use model.fit to train the CNN model, with the same batch_size and number of epochs.
H = model2.fit(train_images,np.array(y_train), validation_data=(test_images , y_test))'''

Epoch 1/40
500/500 [=====] - 4s 7ms/step - loss: 2.0894 - accuracy: 0.2239 - val_loss: 1.8062 - val_accuracy: 0.3672
Epoch 2/40
```

```
500/500 [=====] - 4s 7ms/step - loss: 1.7696 - accuracy: 0.  
3501 - val_loss: 1.6139 - val_accuracy: 0.4131  
Epoch 3/40  
500/500 [=====] - 4s 7ms/step - loss: 1.6299 - accuracy: 0.  
4024 - val_loss: 1.4927 - val_accuracy: 0.4573  
Epoch 4/40  
500/500 [=====] - 4s 7ms/step - loss: 1.5320 - accuracy: 0.  
4385 - val_loss: 1.4038 - val_accuracy: 0.4933  
Epoch 5/40  
500/500 [=====] - 4s 7ms/step - loss: 1.4709 - accuracy: 0.  
4601 - val_loss: 1.3716 - val_accuracy: 0.5079  
Epoch 6/40  
500/500 [=====] - 4s 7ms/step - loss: 1.4003 - accuracy: 0.  
4880 - val_loss: 1.2995 - val_accuracy: 0.5288  
Epoch 7/40  
500/500 [=====] - 4s 7ms/step - loss: 1.3449 - accuracy: 0.  
5112 - val_loss: 1.2440 - val_accuracy: 0.5522  
Epoch 8/40  
500/500 [=====] - 4s 7ms/step - loss: 1.2944 - accuracy: 0.  
5309 - val_loss: 1.2388 - val_accuracy: 0.5566  
Epoch 9/40  
500/500 [=====] - 4s 7ms/step - loss: 1.2482 - accuracy: 0.  
5496 - val_loss: 1.1324 - val_accuracy: 0.5959  
Epoch 10/40  
500/500 [=====] - 4s 7ms/step - loss: 1.1982 - accuracy: 0.  
5724 - val_loss: 1.0941 - val_accuracy: 0.6098  
Epoch 11/40  
500/500 [=====] - 4s 7ms/step - loss: 1.1582 - accuracy: 0.  
5833 - val_loss: 1.0665 - val_accuracy: 0.6172  
Epoch 12/40  
500/500 [=====] - 4s 7ms/step - loss: 1.1245 - accuracy: 0.  
5974 - val_loss: 1.0010 - val_accuracy: 0.6423  
Epoch 13/40  
500/500 [=====] - 4s 7ms/step - loss: 1.0941 - accuracy: 0.  
6096 - val_loss: 0.9900 - val_accuracy: 0.6506  
Epoch 14/40  
500/500 [=====] - 4s 7ms/step - loss: 1.0624 - accuracy: 0.  
6216 - val_loss: 0.9531 - val_accuracy: 0.6643  
Epoch 15/40  
500/500 [=====] - 4s 7ms/step - loss: 1.0301 - accuracy: 0.  
6342 - val_loss: 0.9251 - val_accuracy: 0.6740  
Epoch 16/40  
500/500 [=====] - 4s 7ms/step - loss: 1.0062 - accuracy: 0.  
6416 - val_loss: 0.9171 - val_accuracy: 0.6764  
Epoch 17/40  
500/500 [=====] - 4s 7ms/step - loss: 0.9801 - accuracy: 0.  
6517 - val_loss: 0.9035 - val_accuracy: 0.6776  
Epoch 18/40  
500/500 [=====] - 4s 7ms/step - loss: 0.9582 - accuracy: 0.  
6587 - val_loss: 0.8643 - val_accuracy: 0.6957  
Epoch 19/40  
500/500 [=====] - 4s 7ms/step - loss: 0.9360 - accuracy: 0.  
6685 - val_loss: 0.8467 - val_accuracy: 0.7018  
Epoch 20/40  
500/500 [=====] - 4s 7ms/step - loss: 0.9147 - accuracy: 0.  
6766 - val_loss: 0.8420 - val_accuracy: 0.7035  
Epoch 21/40  
500/500 [=====] - 4s 7ms/step - loss: 0.8927 - accuracy: 0.  
6846 - val_loss: 0.8094 - val_accuracy: 0.7180  
Epoch 22/40  
500/500 [=====] - 4s 7ms/step - loss: 0.8773 - accuracy: 0.  
6901 - val_loss: 0.7961 - val_accuracy: 0.7223  
Epoch 23/40  
500/500 [=====] - 4s 7ms/step - loss: 0.8570 - accuracy: 0.  
6978 - val_loss: 0.7980 - val_accuracy: 0.7239  
Epoch 24/40  
500/500 [=====] - 4s 7ms/step - loss: 0.8423 - accuracy: 0.  
7029 - val_loss: 0.7731 - val_accuracy: 0.7325  
Epoch 25/40
```

```

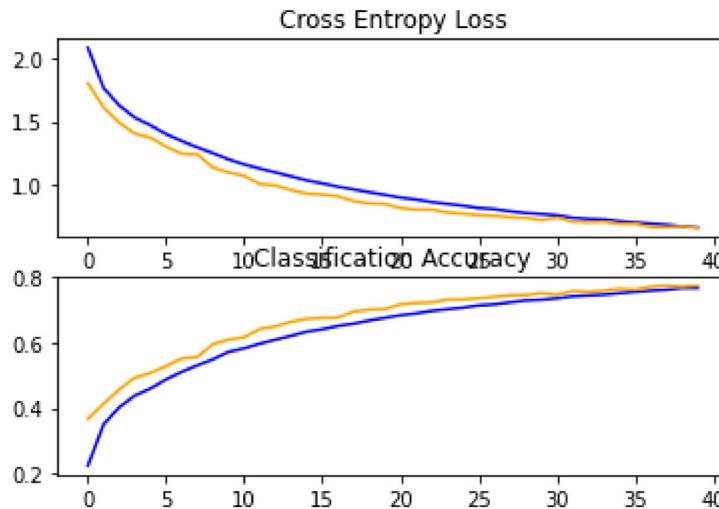
500/500 [=====] - 4s 7ms/step - loss: 0.8280 - accuracy: 0.
7078 - val_loss: 0.7639 - val_accuracy: 0.7332
Epoch 26/40
500/500 [=====] - 4s 7ms/step - loss: 0.8096 - accuracy: 0.
7145 - val_loss: 0.7530 - val_accuracy: 0.7366
Epoch 27/40
500/500 [=====] - 4s 7ms/step - loss: 0.8009 - accuracy: 0.
7178 - val_loss: 0.7472 - val_accuracy: 0.7411
Epoch 28/40
500/500 [=====] - 4s 7ms/step - loss: 0.7834 - accuracy: 0.
7241 - val_loss: 0.7358 - val_accuracy: 0.7456
Epoch 29/40
500/500 [=====] - 4s 7ms/step - loss: 0.7698 - accuracy: 0.
7296 - val_loss: 0.7307 - val_accuracy: 0.7462
Epoch 30/40
500/500 [=====] - 4s 7ms/step - loss: 0.7615 - accuracy: 0.
7319 - val_loss: 0.7132 - val_accuracy: 0.7520
Epoch 31/40
500/500 [=====] - 4s 7ms/step - loss: 0.7519 - accuracy: 0.
7365 - val_loss: 0.7328 - val_accuracy: 0.7470
Epoch 32/40
500/500 [=====] - 4s 7ms/step - loss: 0.7315 - accuracy: 0.
7426 - val_loss: 0.7010 - val_accuracy: 0.7589
Epoch 33/40
500/500 [=====] - 4s 7ms/step - loss: 0.7223 - accuracy: 0.
7453 - val_loss: 0.6956 - val_accuracy: 0.7567
Epoch 34/40
500/500 [=====] - 4s 7ms/step - loss: 0.7188 - accuracy: 0.
7477 - val_loss: 0.6992 - val_accuracy: 0.7601
Epoch 35/40
500/500 [=====] - 4s 7ms/step - loss: 0.7042 - accuracy: 0.
7521 - val_loss: 0.6831 - val_accuracy: 0.7653
Epoch 36/40
500/500 [=====] - 4s 7ms/step - loss: 0.6943 - accuracy: 0.
7560 - val_loss: 0.6846 - val_accuracy: 0.7631
Epoch 37/40
500/500 [=====] - 4s 7ms/step - loss: 0.6842 - accuracy: 0.
7600 - val_loss: 0.6575 - val_accuracy: 0.7718
Epoch 38/40
500/500 [=====] - 4s 7ms/step - loss: 0.6754 - accuracy: 0.
7627 - val_loss: 0.6557 - val_accuracy: 0.7754
Epoch 39/40
500/500 [=====] - 4s 7ms/step - loss: 0.6631 - accuracy: 0.
7681 - val_loss: 0.6635 - val_accuracy: 0.7720
Epoch 40/40
500/500 [=====] - 4s 7ms/step - loss: 0.6537 - accuracy: 0.
7681 - val_loss: 0.6500 - val_accuracy: 0.7752

```

```
In [44]: test_loss, test_acc = model2.evaluate(test_images , y_test)# TODO
print('Test accuracy:', test_acc)
```

```
313/313 [=====] - 1s 3ms/step - loss: 0.6500 - accuracy: 0.
7752
Test accuracy: 0.7752000093460083
```

```
In [45]: summarize_diagnostics(H)
```

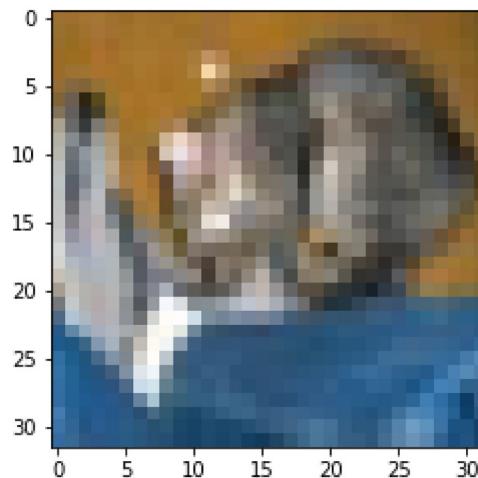


In [72]:

```
for i in range(10):
    predictions = model2.predict(np.array([test_images[i]]))
    prediction = np.argmax(predictions[0])# TODO
    print(f"Class image : {prediction +1}")
    print("Label of this digit is:", y_test[i])
    plt.imshow(test_images[i,:,:], cmap=plt.cm.binary)
    plt.show()
```

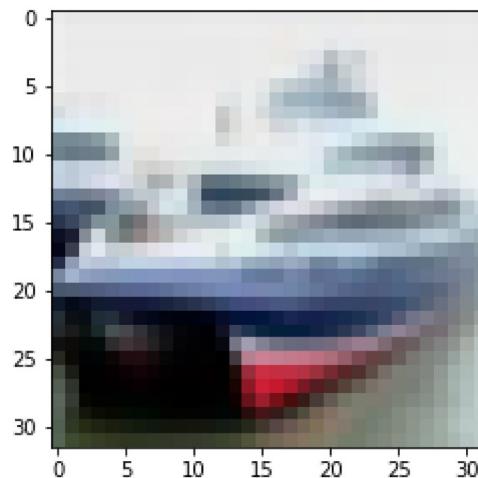
Class image : 4

Label of this digit is: [0 0 0 1 0 0 0 0 0 0]



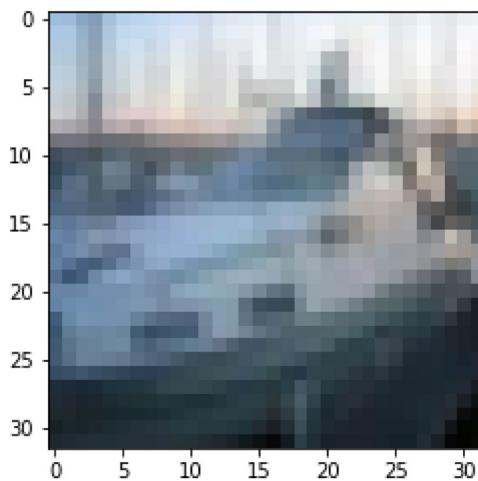
Class image : 9

Label of this digit is: [0 0 0 0 0 0 0 0 1 0]



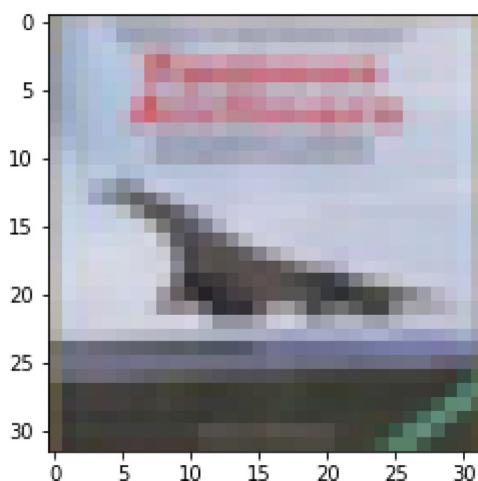
Class image : 9

Label of this digit is: [0 0 0 0 0 0 0 0 1 0]



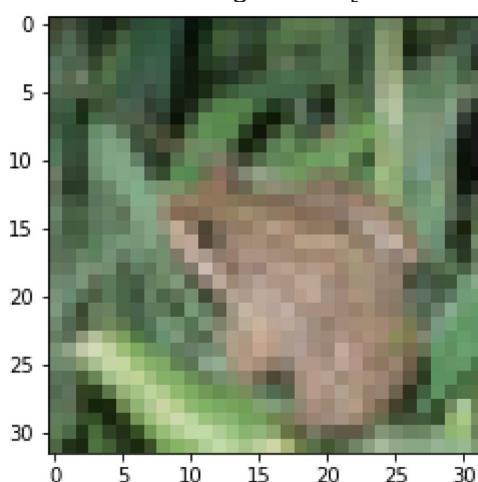
Class image : 1

Label of this digit is: [1 0 0 0 0 0 0 0 0 0]



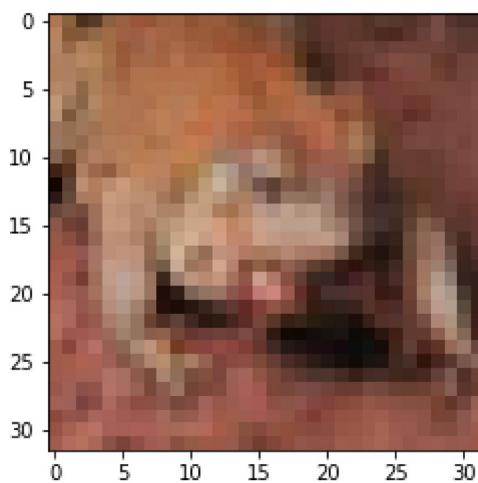
Class image : 7

Label of this digit is: [0 0 0 0 0 0 1 0 0 0]



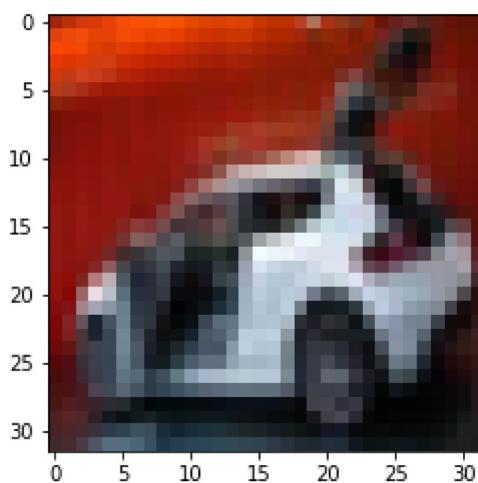
Class image : 7

Label of this digit is: [0 0 0 0 0 0 1 0 0 0]



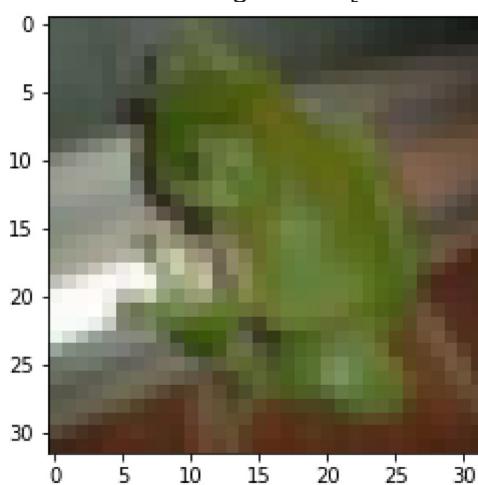
Class image : 2

Label of this digit is: [0 1 0 0 0 0 0 0 0 0]



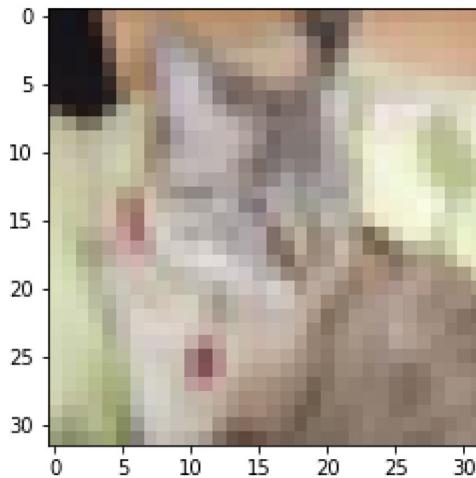
Class image : 7

Label of this digit is: [0 0 0 0 0 0 1 0 0 0]

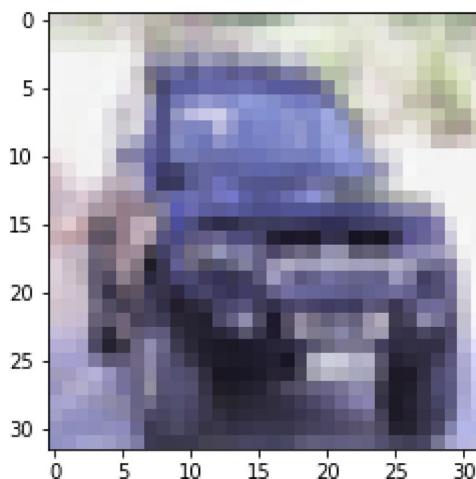


Class image : 4

Label of this digit is: [0 0 0 1 0 0 0 0 0 0]



Class image : 2
Label of this digit is: [0 1 0 0 0 0 0 0 0 0]



```
In [62]: # predictions = model2.predict(np.array([test_images[2]]))
```

```
In [63]: # predictions
```

```
Out[63]: array([[8.0983527e-03, 6.2280190e-03, 9.2255854e-05, 1.2325000e-03,
   3.0086128e-05, 2.9635085e-05, 1.0467141e-05, 7.0712020e-05,
   9.7792393e-01, 6.2839459e-03]], dtype=float32)
```

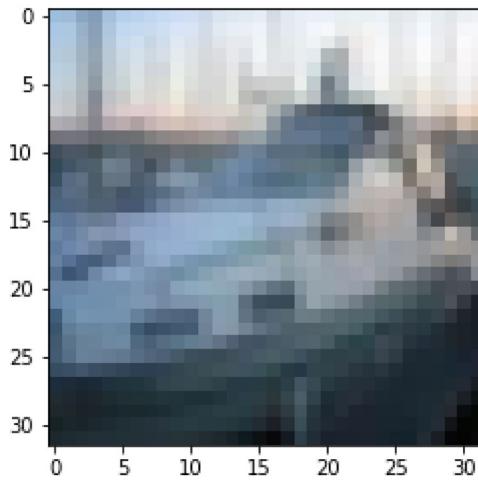
```
In [64]: # prediction = np.argmax(predictions[0])# TODO
# print(f"Class image : {prediction +1}")
```

Class image : 9

```
In [65]: # print("Label of this digit is:", y_test[2])
# plt.imshow(test_images[2,:,:], cmap=plt.cm.binary)
```

Label of this digit is: [0 0 0 0 0 0 0 1 0]

```
Out[65]: <matplotlib.image.AxesImage at 0x7f3625ee0890>
```



```
In [74]: from tensorflow.python.client import device_lib  
device_lib.list_local_devices()
```

```
Out[74]: [name: "/device:CPU:0"  
device_type: "CPU"  
memory_limit: 268435456  
locality {  
}  
incarnation: 11052314551703897330, name: "/device:GPU:0"  
device_type: "GPU"  
memory_limit: 16183459840  
locality {  
    bus_id: 1  
    links {  
    }  
}  
incarnation: 10008767861492315563  
physical_device_desc: "device: 0, name: Tesla P100-PCIE-16GB, pci bus id: 0000:00:0  
4.0, compute capability: 6.0"]
```